

AWS Databases

AWS Database Options

RDS

Non-
RDS

In-Memory

Data
Warehousing

Custom

Running Databases on AWS:

http://aws.amazon.com/running_databases/

	SQL	NoSQL
Data Storage	Rows and Columns	Key-Value
Schemas	Fixed	Dynamic
Querying	Using SQL	Focused on collection of documents
Scalability	Vertical	Horizontal

SQL

ISBN	Title	Author	Format
9182932465265	Cloud Computing Concepts	Wilson, Joe	Paperback
3142536475869	The Database Guru	Gomez, Maria	eBook

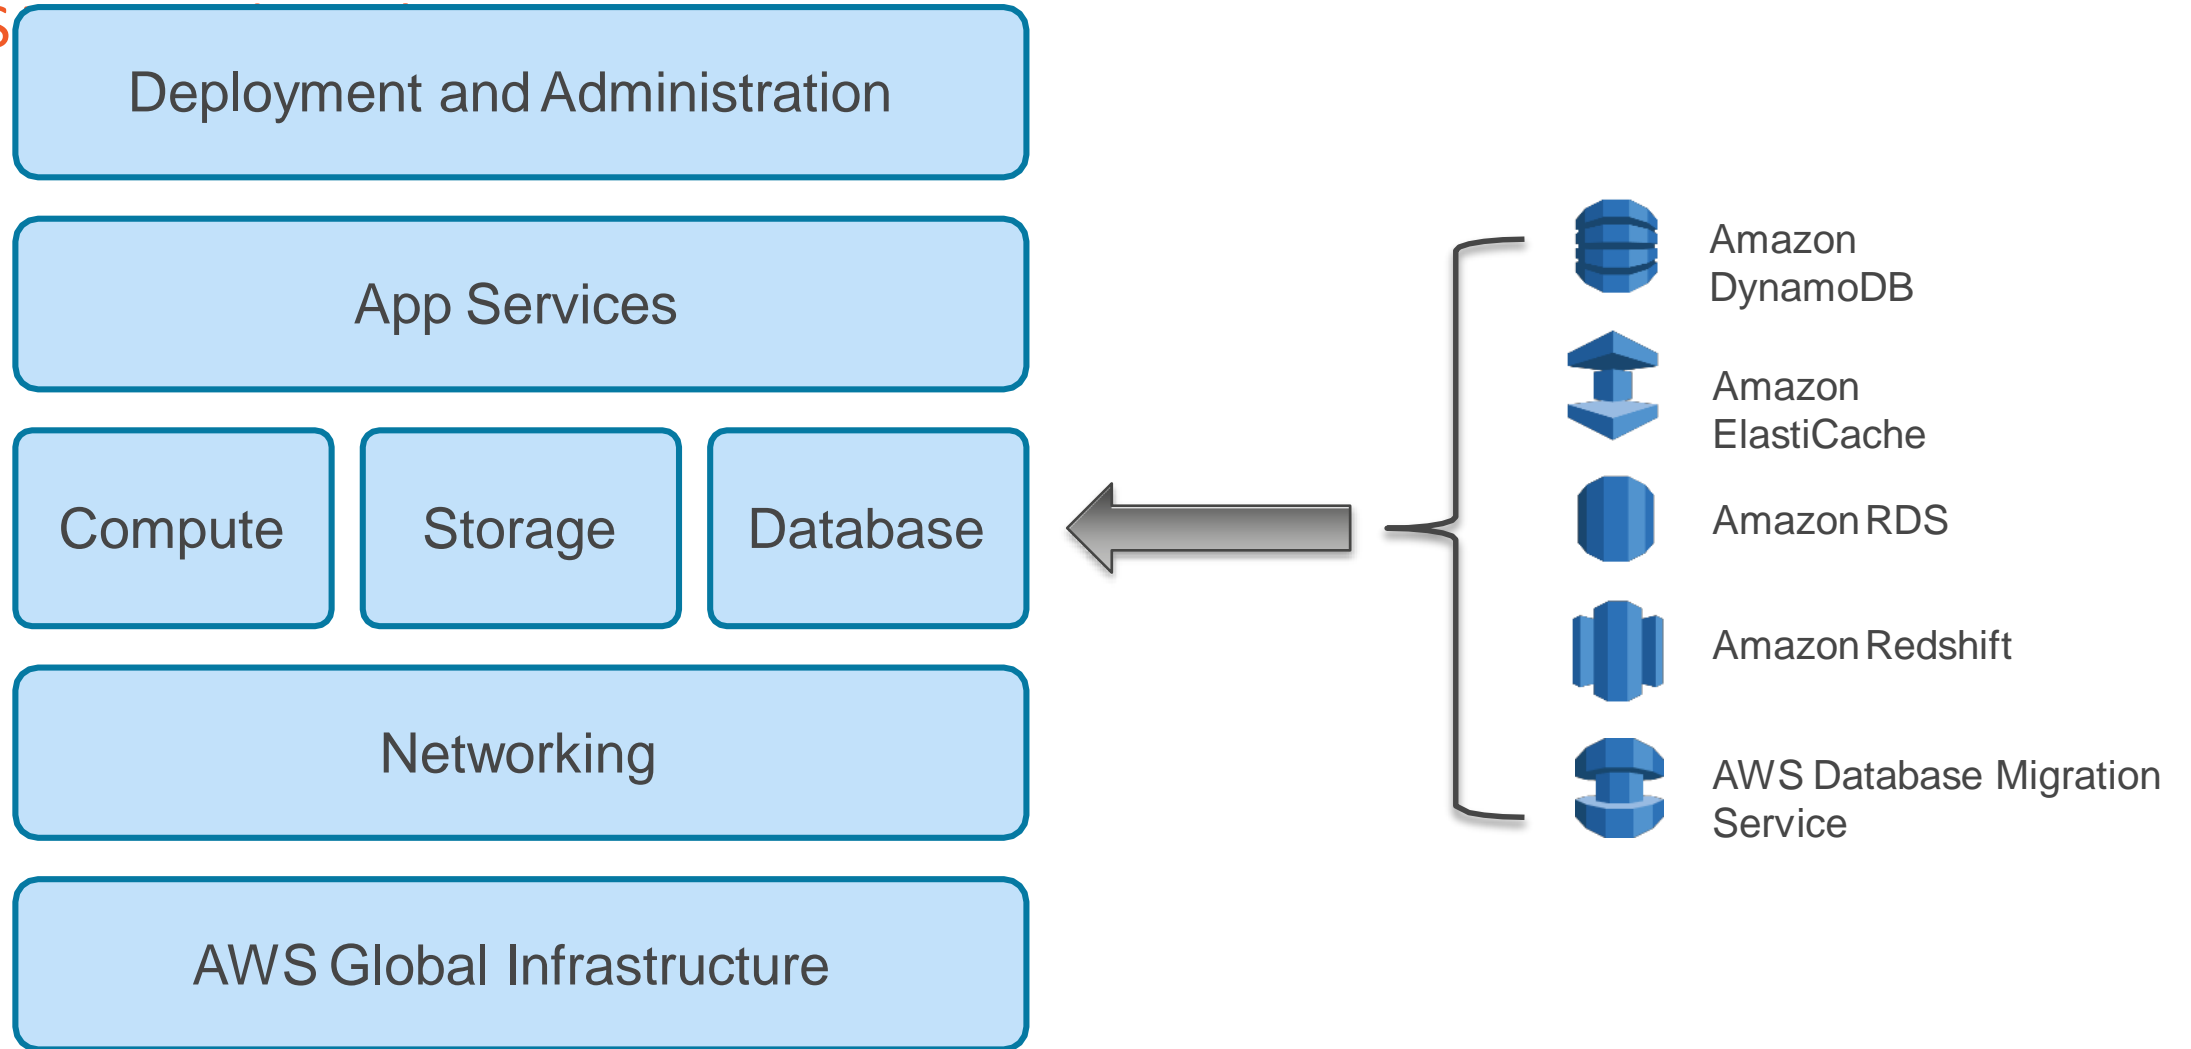
NoSQL

```
{
  ISBN: 9182932465265,
  Title: "Cloud Computing Concepts",
  Author: "Wilson, Joe",
  Format: "Paperback"
}
```

Data Storage Considerations

- No one size fits all.
- Analyze your data requirements by considering:
 - Data formats
 - Data size
 - Query frequency
 - Data access speed
 - Data retention period

AWS



Amazon RDS Options

MySQL

Oracle

Microsoft SQL

PostgreSQL

Aurora

Non-RDS

SimpleDB

DynamoDB

MongoDB

Couchbase

In - Memory

Elastic Cache4

Datebase warehouse

Amazon Redshift

Amazon Relational Database Service (RDS)



Amazon
RDS

- Cost-efficient and **resizable capacity**
- Manages time-consuming **database administration** tasks
- Access to the full capabilities of **Amazon Aurora, MySQL, MariaDB, Microsoft SQL Server, Oracle, and PostgreSQL** databases



Amazon RDS


- Simple and **fast to deploy**
- Manages common database administrative tasks
- **Compatible** with your applications
- Fast, predictable performance
- Simple and **fast to scale**
- Secure
- Cost-effective





DB Instances

- DB Instances are the basic building blocks of Amazon RDS.
- They are an **isolated database environment** in the cloud.
- They can **contain multiple user-created databases**.



How Amazon RDS Backups Work

Automatic Backups:

- Restore your database to a point in time.
- Are enabled by default.
- Let you choose a retention period up to 35 days.



Manual Snapshots:

- Let you build a new database instance from a snapshot.
- Are initiated by the user.
- Persist until the user deletes them.
- Are stored in Amazon S3.



Cross-Region Snapshots

- Are a **copy** of a **database** snapshot stored in a **different AWS Region**.
- Provide a backup for disaster **recovery**.
- Can be used as a **base** for **migration** to a different region.






Amazon RDS Security

- Run your DB instance in an **Amazon VPC**.
- Use **IAM policies** to grant access to Amazon RDS resources.
- Use **security groups**.
- Use Secure Socket Layer (**SSL**) connections with DB instances (Amazon Aurora, Oracle, MySQL, MariaDB, PostgreSQL, Microsoft SQL Server).
- Use Amazon RDS **encryption** to secure your RDS instances and snapshots at rest.
- Use network encryption and transparent data encryption (**TDE**) with Oracle DB and Microsoft SQL Server instances.
- Use the security features of your DB engine to **control access** to your DB instance.

A Simple Application Architecture






Multi-AZ RDS Deployment

- With **Multi-AZ** operation, your database is **synchronously replicated to another Availability Zone** in the same AWS Region.
- **Failover** to the standby **automatically** occurs in case of master database failure.
- Planned maintenance is applied first to standby databases.

A Resilient, Durable Application Architecture



Elastic Load Balancing
load balancer instance

Application, in Amazon
EC2 instances

Amazon RDS database instances:
Master and Multi-AZ standby

DB snapshots in
Amazon S3



Amazon RDS Best Practices

- **Monitor** your memory, CPU, and storage usage.
- Use **Multi-AZ** deployments to automatically provision and maintain a synchronous standby in a different Availability Zone.
- Enable **automatic backups**.
- Set the **backup window** to occur during the daily low in WriteIOPS.
- To increase the I/O capacity of a DB instance:
 - Migrate to a DB instance class with high I/O capacity.
 - Convert from standard storage to provisioned IOPS storage and use a DB instance class optimized for **provisioned IOPS**.
 - Provision additional throughput capacity (if using provisioned IOPS storage).
- **Test** failover for your DB instance.

Amazon DynamoDB

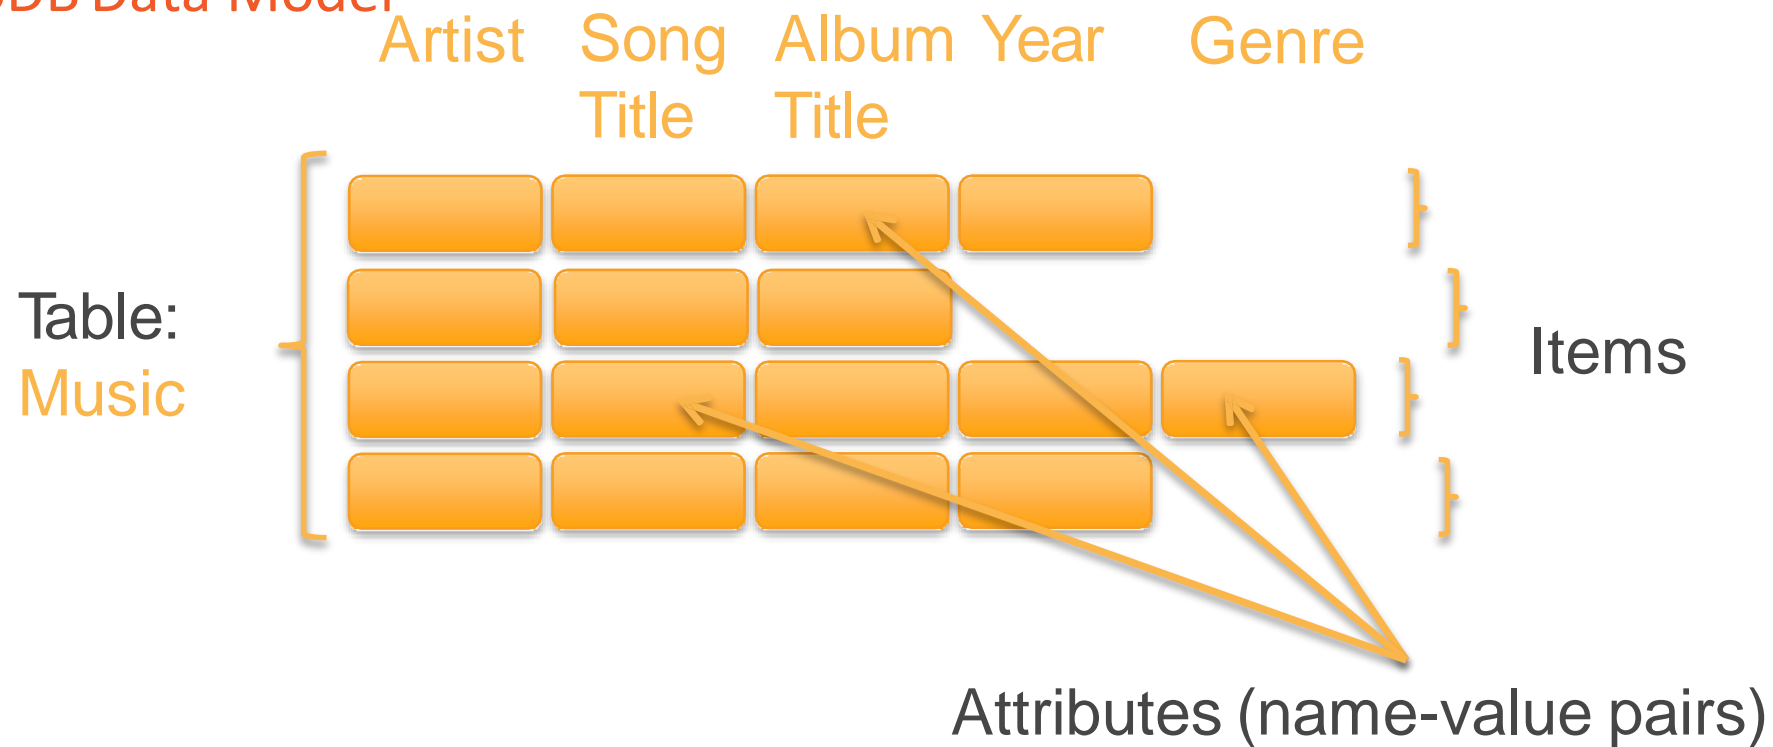


Amazon
DynamoDB

- Allows you to store any amount of data with **no limits**.
- Provides fast, predictable performance using **SSDs**.
- Allows you to easily provision and change the **request capacity** needed for each table.
- ~~Is a **fully managed, NoSQL** database service.~~



DynamoDB Data Model





Primary Keys

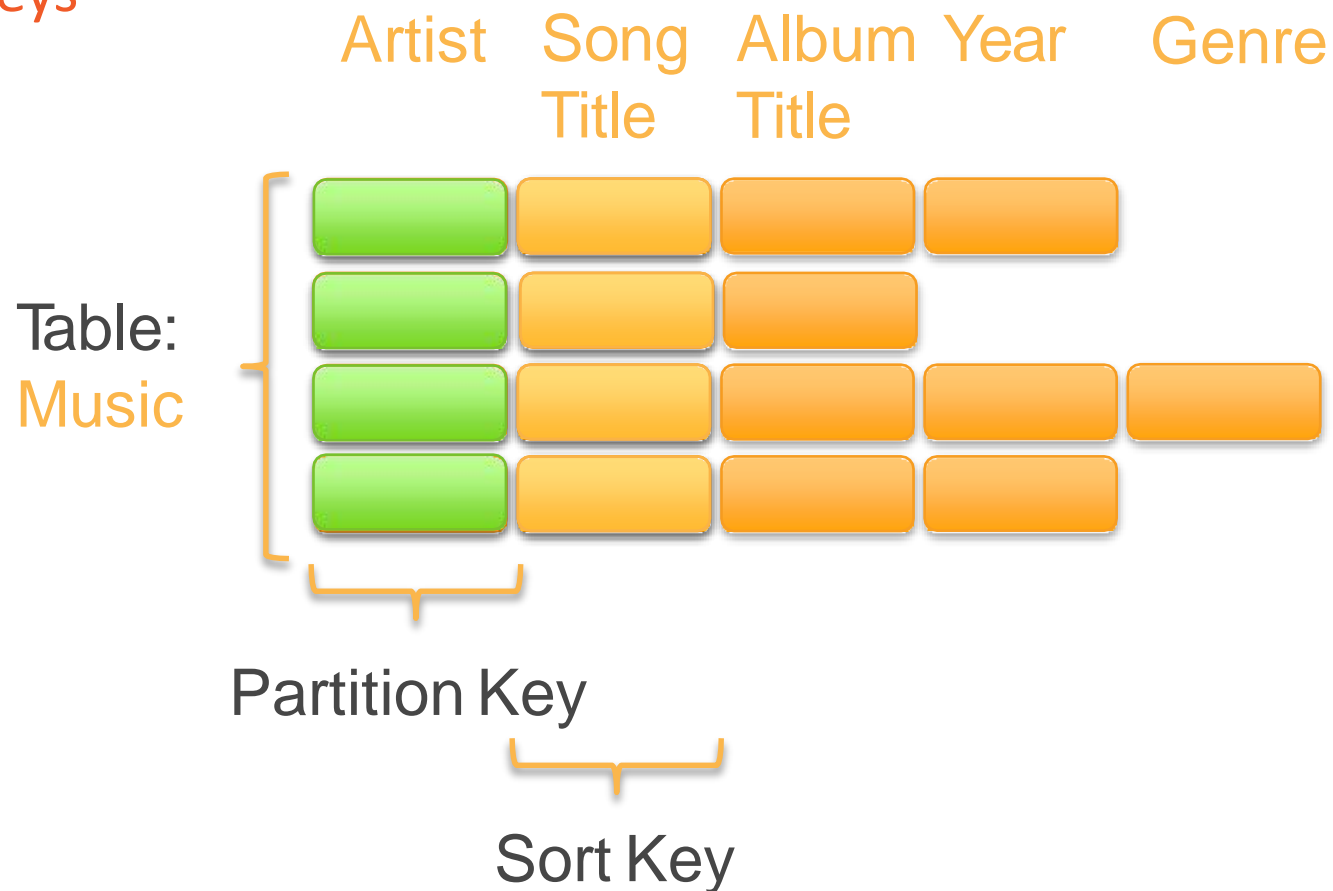


Table: Music
Partition Key: Artist
Sort Key: Song Title

(DynamoDB maintains a sorted index for both keys)



Provisioned Throughput

You specify how much **provisioned throughput capacity** you need for reads and writes.

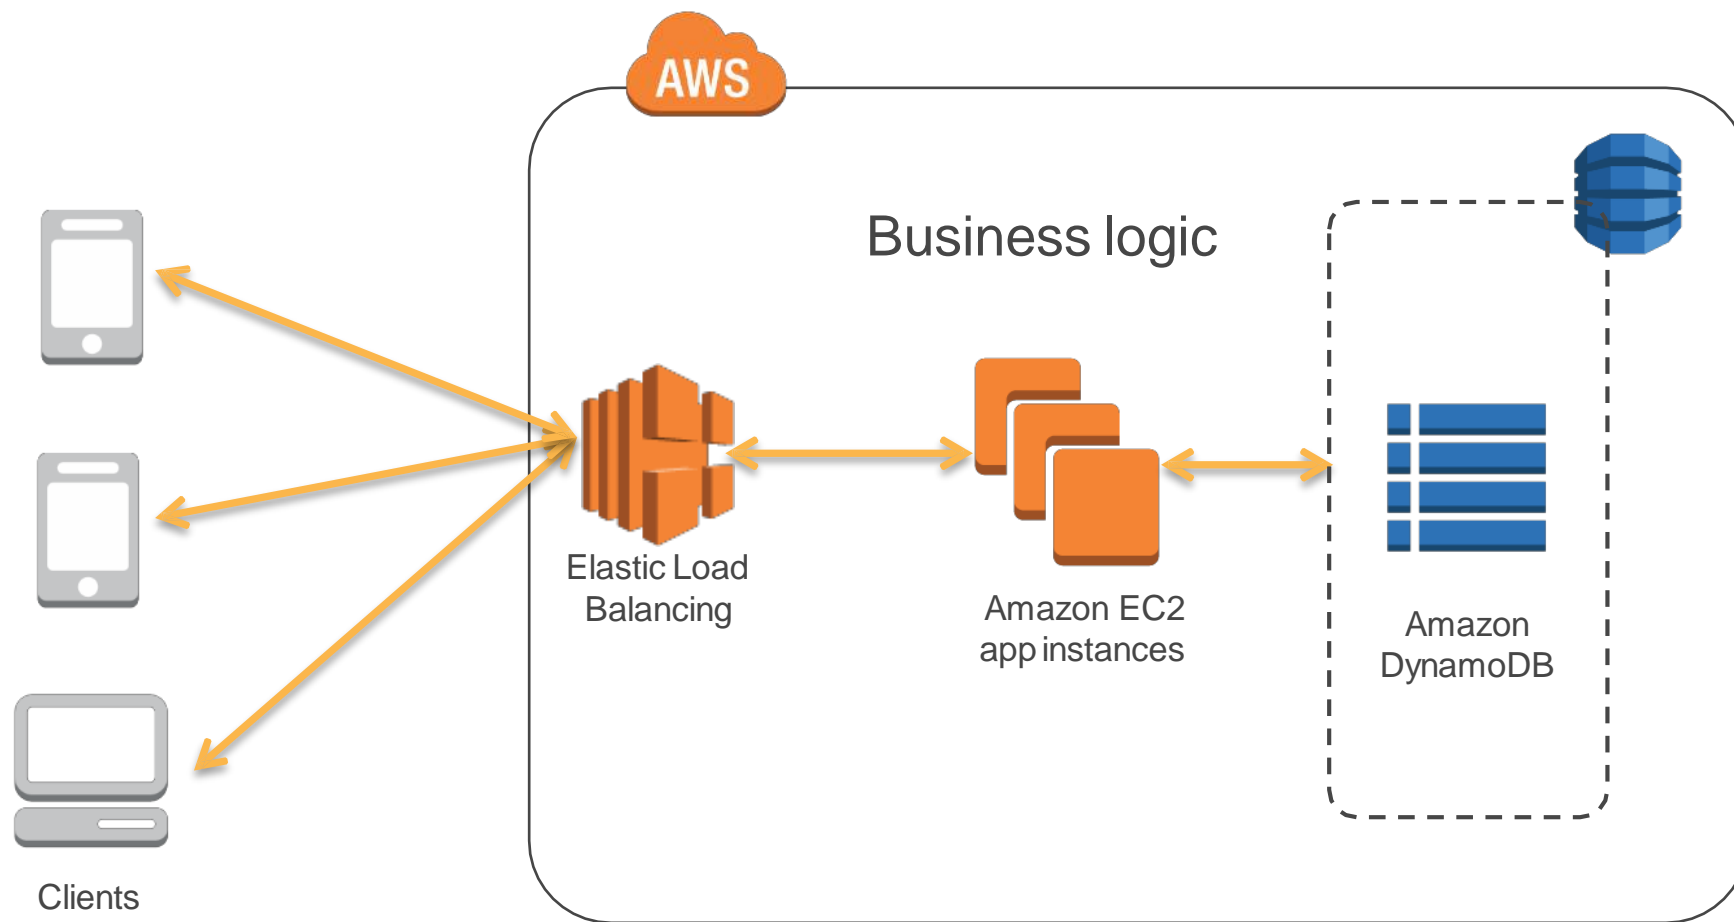
Amazon DynamoDB allocates the necessary machine resources to meet your needs.








Supported Operations

- **Query:**
 - Query a table using the partition key and an optional sort key filter.
 - If the table has a secondary index, query using its key.
 - It is the **most efficient way to retrieve items** from a table or secondary index.
- **Scan:**
 - You can scan a table or secondary index.
 - Scan reads every item – **slower than querying**.
- You can use conditional expressions in both Query and Scan operations.

Simple Application Architecture



Factors	Relational (Amazon RDS) 	NoSQL (Amazon DynamoDB) 
Application Type	<ul style="list-style-type: none"> Existing database apps Business process—centric apps 	<ul style="list-style-type: none"> New web-scale applications Large number of small writes and reads
Application Characteristics	<ul style="list-style-type: none"> Relational data models, transactions Complex queries, joins, and updates 	<ul style="list-style-type: none"> Simple data models, transactions Range queries, simple updates
Scaling	Application or DBA—architected (clustering, partitions, sharding)	Seamless, on-demand scaling based on application requirements
QoS	<ul style="list-style-type: none"> Performance—depends on data model, indexing, query, and storage optimization Reliability and availability Durability 	<ul style="list-style-type: none"> Performance—Automatically optimized by the system Reliability and availability Durability

If You Need	Consider Using
A relational database service with minimal administration	Amazon RDS <ul style="list-style-type: none"> • Choice of Amazon Aurora, MySQL, MariaDB, Microsoft SQL Server, Oracle, or PostgreSQL database engines • Scale compute and storage • Multi-AZ availability 
A fast, highly scalable NoSQL database service	Amazon DynamoDB <ul style="list-style-type: none"> • Extremely fast performance • Seamless scalability and reliability • Low cost 
A database you can manage on your own	Your choice of AMIs on Amazon EC2 and Amazon EBS that provide scale compute and storage, complete control over instances, and more. 

Knowledge Check

Q: What are the basic building blocks of Amazon Relational Database Service (RDS)?

DB Instances

True or False: Amazon DynamoDB allows you to store any amount of data with no limits.

True

True or False: Scan is the most efficient way to retrieve items from a DynamoDB table.

False

Q: You are creating a resilient, durable application using Amazon RDS. In addition to Amazon RDS's automatic backups, what feature should you use to ensure that your backups are durable retained?

Manual Snapshots

Understanding AWS Relational Database Services (RDS)



RDS

Characteristics

Database engine managed by AWS

MySQL, Oracle, Microsoft SQL, PostgreSQL, MariaDB, or Amazon Aurora

Multi-AZ deployment options

On-demand and reserved instance pricing

Magnetic, GP-SSD, or PIOPS

Oracle and Microsoft SQL licensing:

- Included Licenses
- Bring your own licensing

Automated or manual backups

RDS Automated Backups

Continuously tracks changes
and backs up your DB

Volume snapshot of your entire DB
instance, not just DBs

One day of backups retained by default
but can be configured up to 35 days

Backup retention period
defined during configuration

When you delete an RDS instance, all
automated snapshots are deleted

- Manual snapshots are preserved

RDS Automated Backups

Automated backups occur daily during a 30 minute configurable backup window

Automated backups are preserved for a configurable number of days (retention period)

RDS Restore

You cannot restore from a DB snapshot to an existing DB instance

- A new DB instance is created when you restore

Only default DB parameters and security groups are restored

- You must manually associate all other DB parameters and SGs

RDS combines daily backups in conjunction with transaction logs to restore the DB Instance to any point during the retention period

Up to the last five minutes

Multi-AZ Failover

Multi-AZ RDS deployment designed for HA


Synchronous replica in secondary AZ

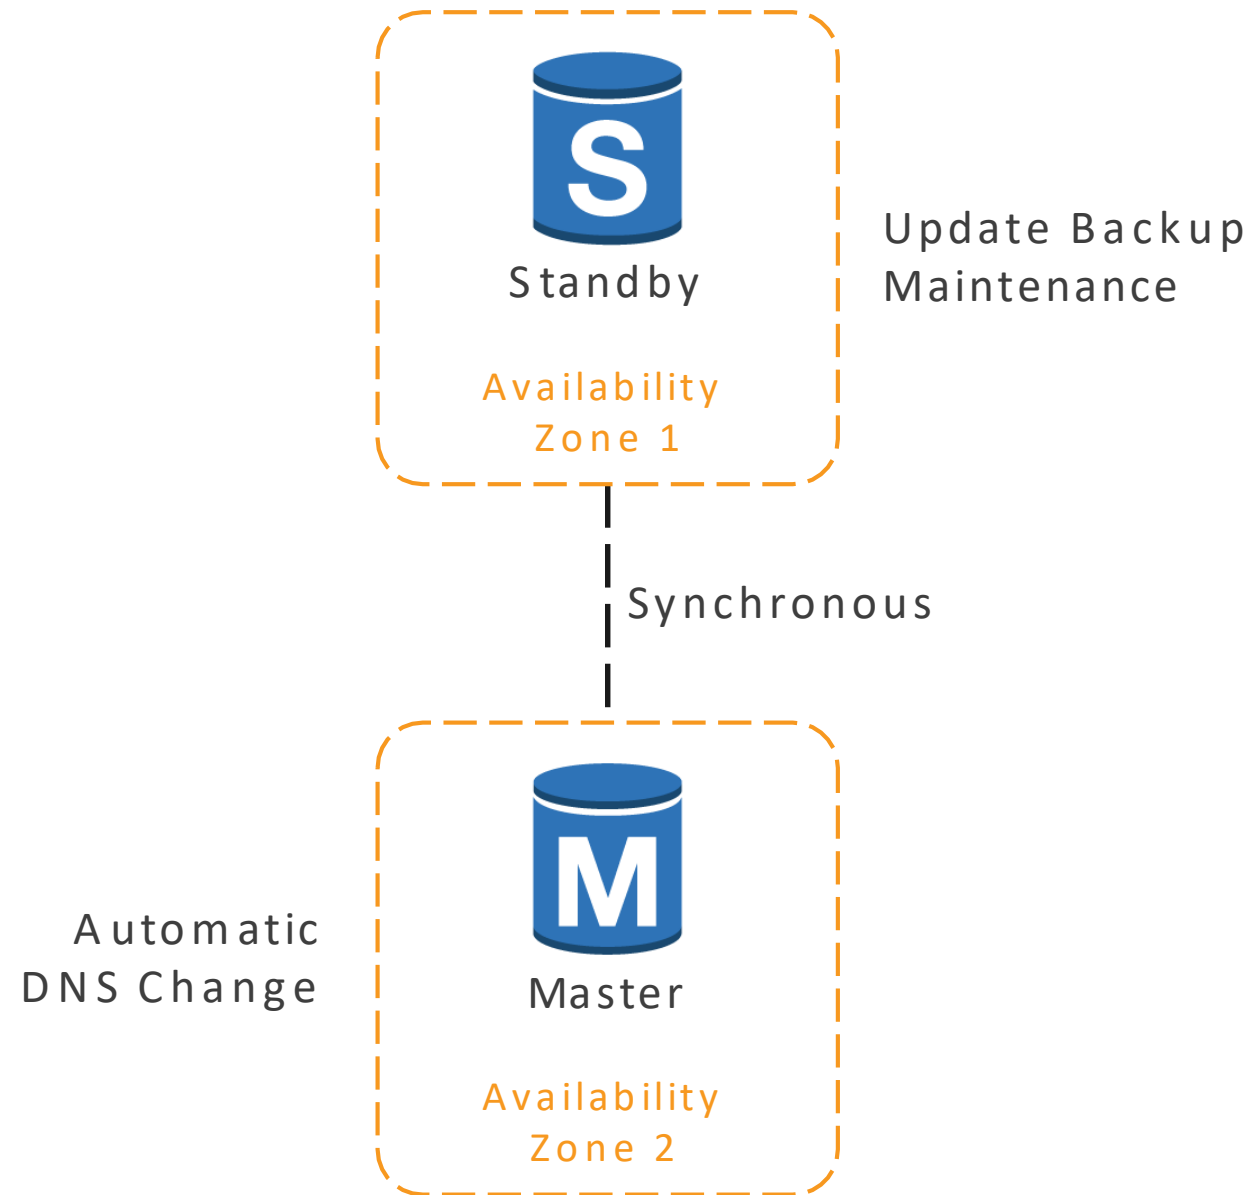
Standby replica RDS instance is invisible

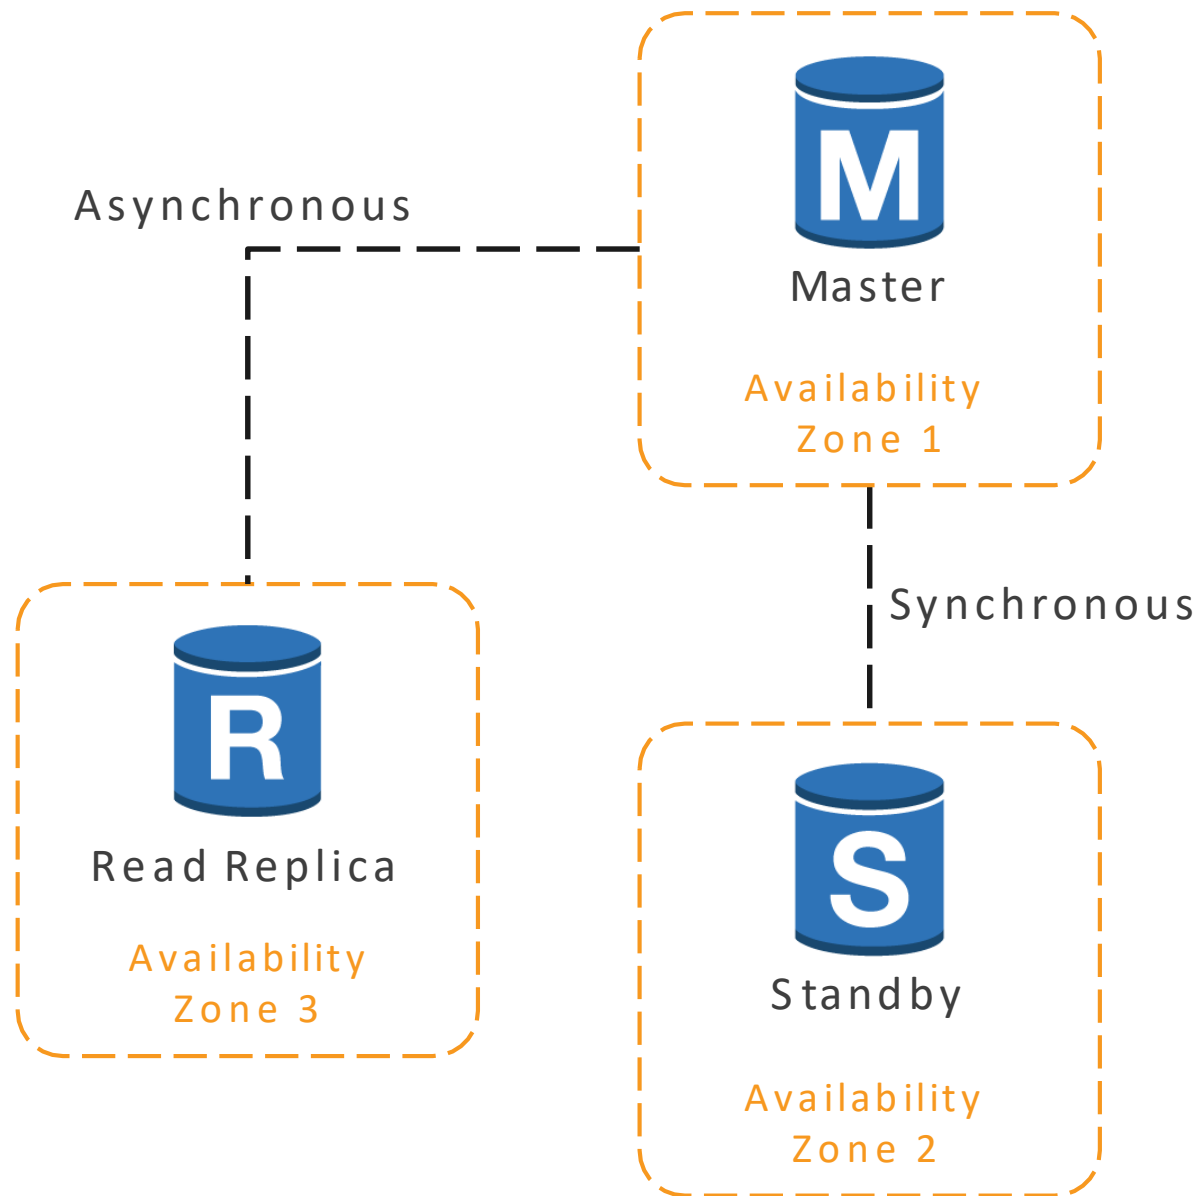
DB snapshots always taken
against standby instance

AWS automatically adjusts DNS record
when needed

Multi-AZ is different from a
RDS read replica







RDS Read Replicas

Read replicas designed for workload sharing / offloading

Created from a snapshot of the master instance

Asynchronous replication / Read-only connections

Read-only disaster recovery

RDS Reserved Instances

DB engine

DB instance class

Deployment type

License model

Region


RDS Reserved Instances

Move between Azs in the same Region

Are available for Multi-AZ deployments

Can be applied to Read Replicas provided the DB Instance class and Region are the same

Demo



Creating an RDS instance

Windows Integrated Authentication

Choose one of the AWS offered directory services

Establish a trust relationship

Windows integrated authentication only works with a domain created using AWS directory service

Alternatively, you can use SQL authentication

Summary



RDS characteristics

Multi-AZ failover

RDS read replicas

RDS reserved instances

Understanding Amazon DynamoDB and Redshift



Amazon DynamoDB

Fully managed, highly available and scalable NoSQL

Automatically and synchronously replicates data across three Availability Zones

SSDs and limiting indexing on attributes provides high throughput and low latency

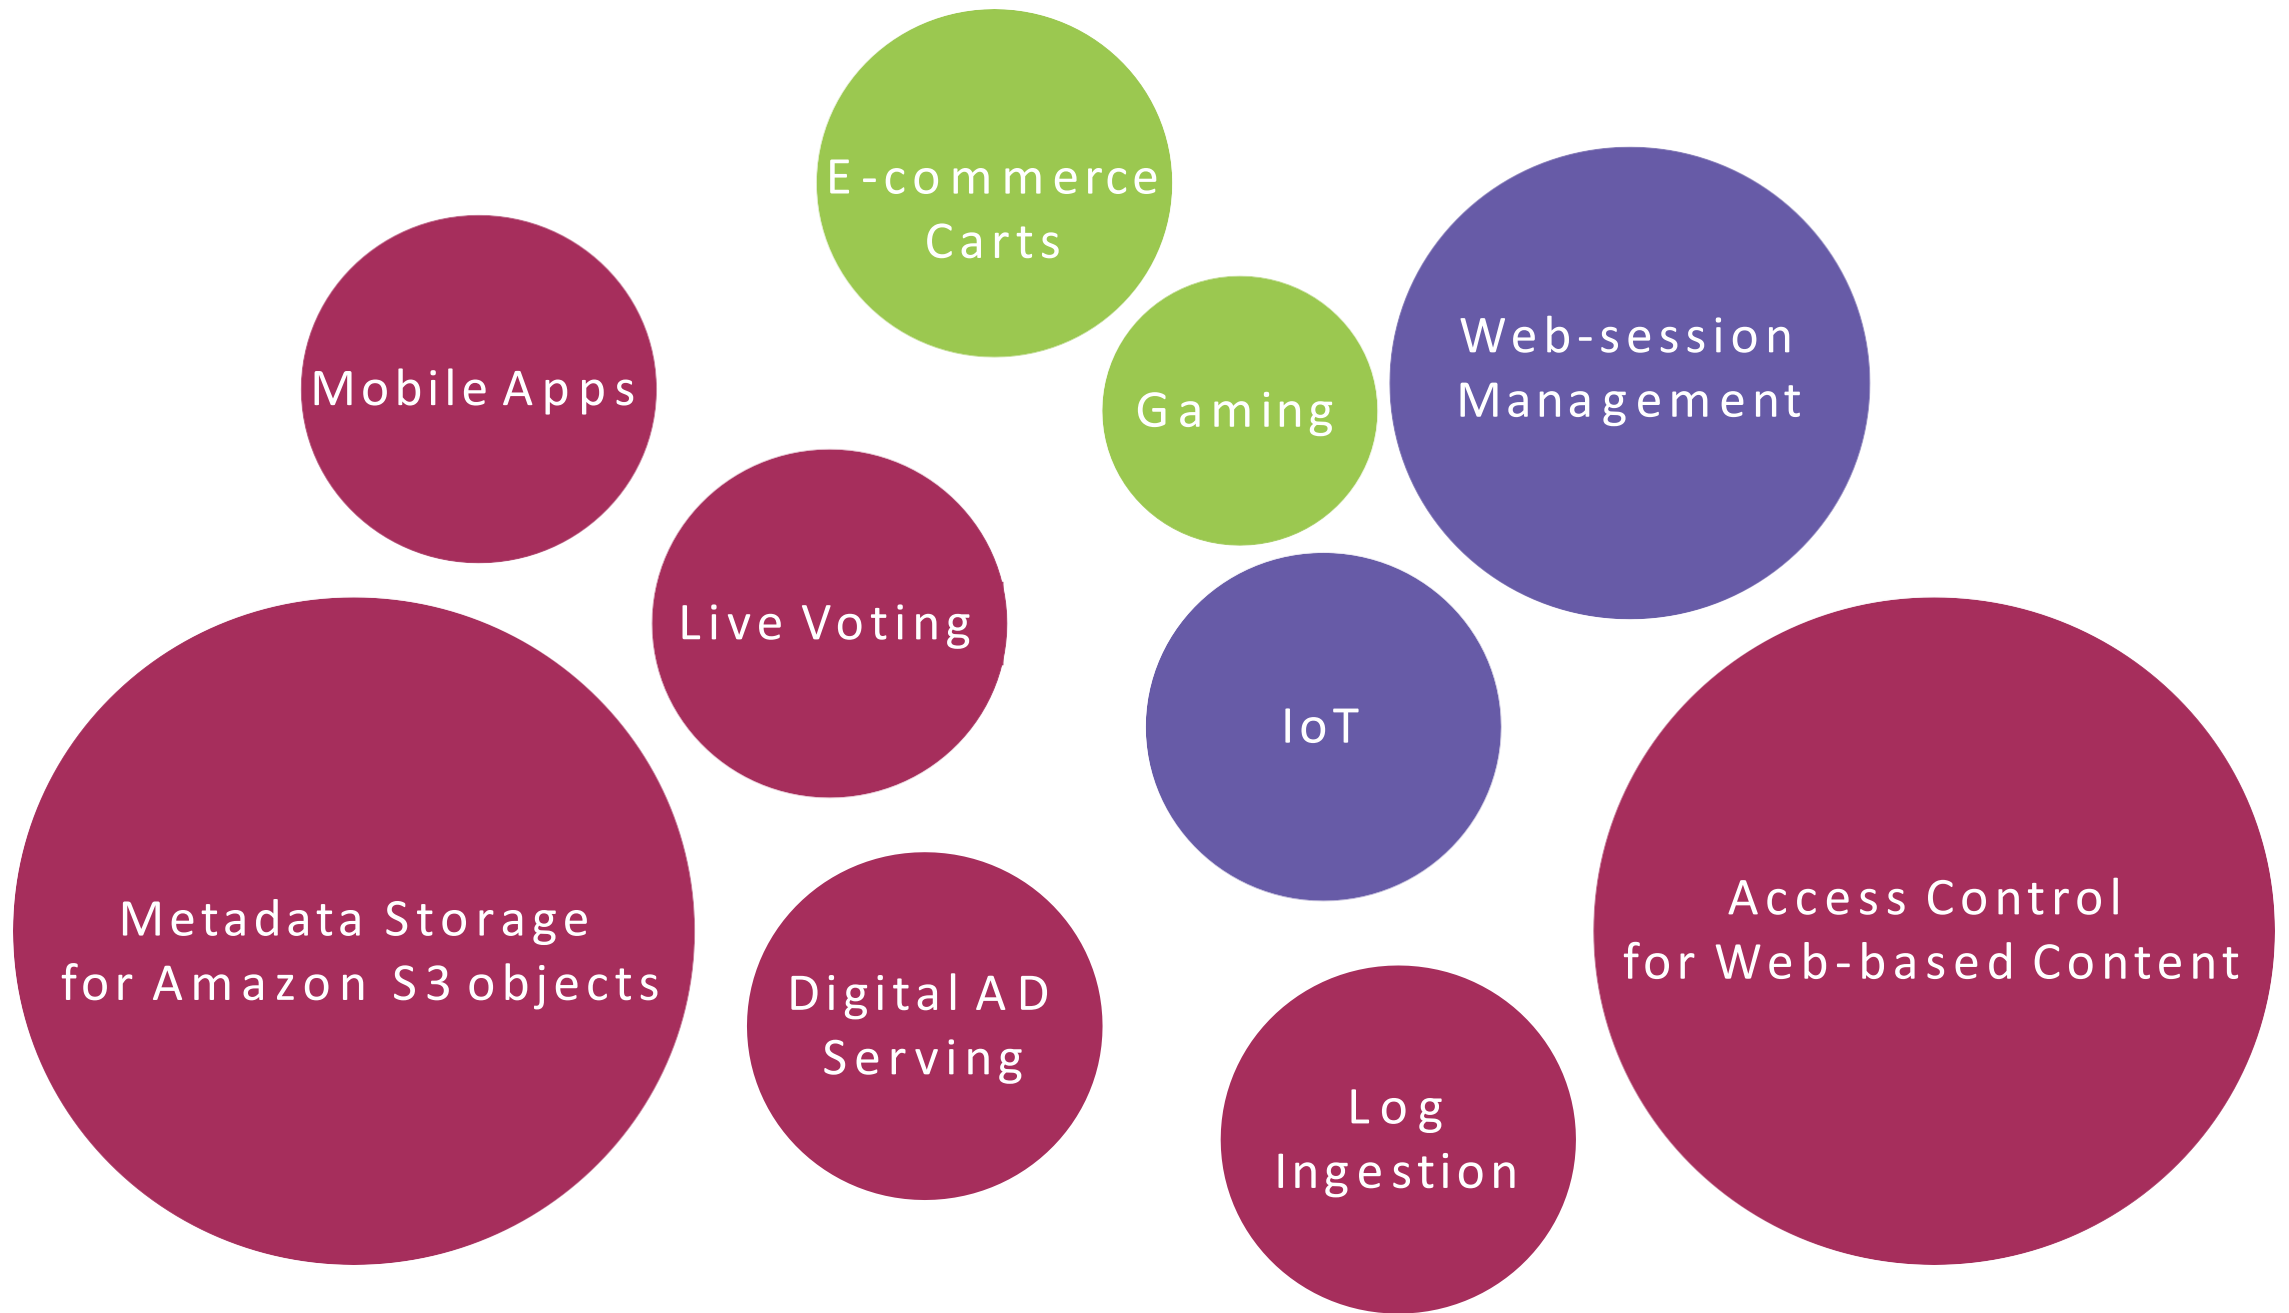
ElastiCache can be used in front of DynamoDB

- Offload high amounts of reads for non- frequently changed data

Amazon DynamoDB

Ideal for existing or new applications that need:

- Flexible NoSQL database with low read and write latencies
- Ability to scale storage and throughput up or down as needed without code changes or downtime



Non-ideal DynamoDB Scenarios

Prewritten application
tied to a traditional
relational database

Joins and/or complex
transactions

BLOB data

Large data with low I/O rate

DynamoDB Integration

Amazon Elastic
MapReduce

Amazon Redshift

Amazon Data
Pipeline

Amazon S3

Management
Console and APIs

Amazon DynamoDB

Stores structured data in tables, indexed by a primary key

Tables are a collection of items and Items are made up of attributes (columns)

Primary key can be:

- Single-attribute hash key
- Composite hash-range key

DynamoDB Features

Secondary indexes



Streams

Cross-region
replication

Triggers

Schema-less

Two Ways to Search



ElastiCache

Open-source in-memory caching engines

- Memcached
 - Widely adopted memory object caching system
- Redis
 - Popular open-source in-memory key-value store
 - Supports data structures such as sorted sets and lists

Master / Slave replication and Multi-AZ

- Can be used to achieve cross AZ redundancy

Memcached vs. Redis

	Memcached	Redis
Cache to offload DB	√	√
Multithreaded performance	√	X
Horizontal scaling	√	X
Multi-AZ	X	√
Backup and restore	X	√
Pub/sub functionality	X	√
Sorting and ranking	X	√
Advanced data types	X	√
Persistence	X	√

Suggested Reading

Performance at Scale with Amazon ElastiCache:

<https://d0.awsstatic.com/whitepapers/performance-at-scale-with-amazon-elasticache.pdf>



Amazon Redshift

Fast and fully managed petabyte-scale relational data warehouse service

Analyze all your data using your existing business intelligence tools

HDD and SSD Platforms

Starts at \$0.25/hour

Scale to \$1,000/TB/Year

Amazon Redshift Architecture

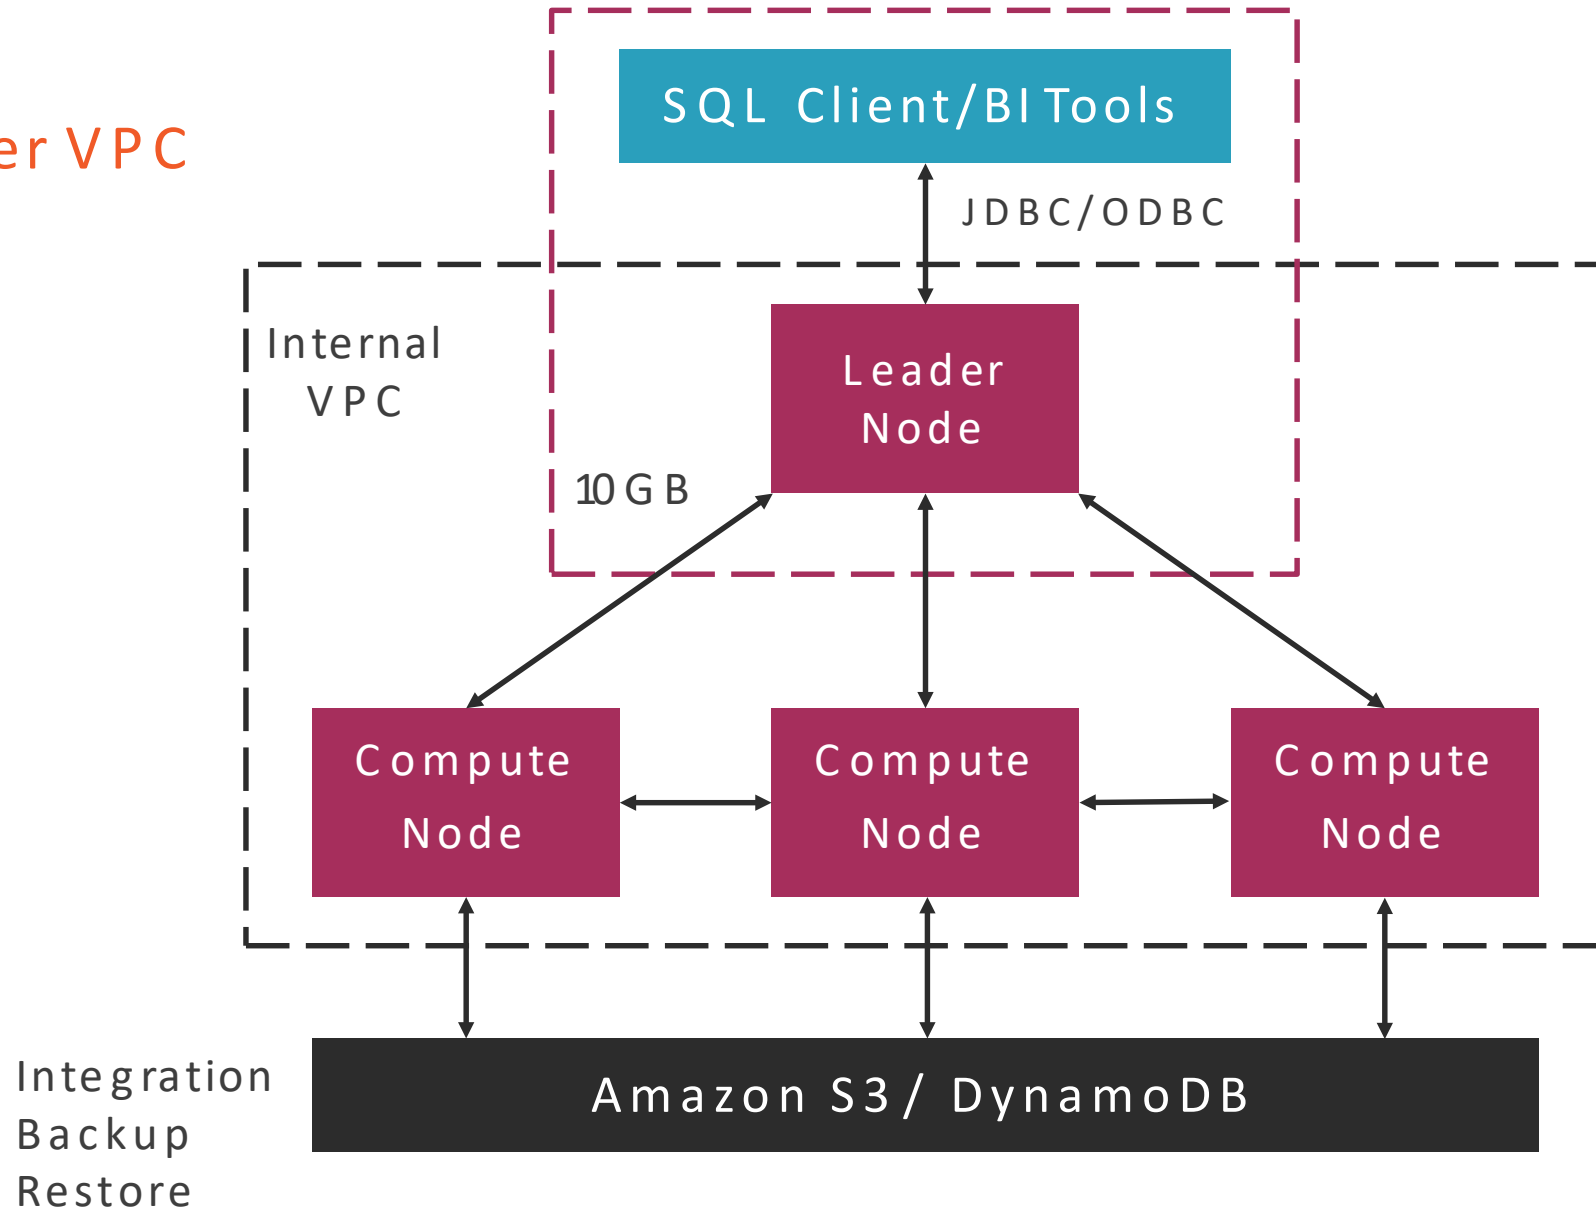
Leader Node

- Simple SQL end point
- Stores metadata
- Optimizes query plan
- Coordinates query execution

Compute Nodes

- Local columnar storage
- Parallel/distributed execution of all queries, loads, backups, restores, resizes

Customer VPC



Backup and Fault Tolerance

Continuous/incremental backups

- Multiple copies within cluster
- Continuous and incremental backups to S3
- Continuous and incremental backups across regions
- Streaming restore

Backup and Fault Tolerance

Fault tolerance

- Disk failures
- Node failures
- Network failures
- Availability Zone/Region level disasters

Security

Load encrypted from S3

SSL to secure data in transit
Amazon VPC for network

isolation Encryption to secure
data at rest

Audit logging and AWS
CloudTrail integration

SOC 1/2/3, PCI-DSS, FedRAMP,
BAA

Amazon DynamoDB

Amazon Redshift